

ARTIFICIAL NEURAL NETWORKS II

Prof. Alexander Huth

3.31.2020

COURSE ADMIN / ZOOM

- * Welcome to the new zoom era
- * If you're watching this, you've found the lecture links posted on Canvas
- * This (and future) lectures will be **recorded** and shared with the class via Canvas
 - * If you don't want your face recorded, please disable your video!
 - * Please don't share the recordings outside of class!
- * And please don't share the zoom link outside of class!

COURSE ADMIN / ZOOM

- * Current plan for zoom etiquette:
 - * Please **keep yourselves muted** except when you are actively speaking
 - * If you want to ask a question, **use the “Raise hand” button** and I’ll call on you
 - * If I miss your raised hand (apologies in advance), toggle it a few times, or just un-mute yourself and ask
 - * I probably won’t be able to keep track of chat messages while lecturing & screen-sharing, so **please avoid using the chat feature**

COURSE ADMIN / ZOOM

- * Office hours will also happen on zoom, Tuesdays & Wednesdays, 10:30am-12:00pm
- * I'll start a zoom meeting & post the link on canvas around when my office hours start

COURSE ADMIN / HW

- * Homework 1 is not graded yet, but will hopefully be returned to you in ~1 week
- * Homework 2 covers linearized model comparison and variance partitioning. It is posted **today** & will be due in 2 weeks (**April 14**)

COURSE ADMIN / PROJECT

- * Goal is to apply or explore something / anything we've talked about in this class
- * could be using real data (e.g. fit some kind of model to a neural dataset)
- * could be theory/methods (e.g. find a better way to do something)

COURSE ADMIN / PROJECT

- * Since there are so many people in this class, I will require you to work on your projects in groups of 2-4
- * If you can't find a group, please contact me or come to office hours to discuss

COURSE ADMIN / PROJECT

- * Proposal due next Thursday (April 9):
 - * ~1-2 paragraphs describing what you plan to do & who is in your group. Email to huth@cs.utexas.edu before class
- * Writeup (3-4 pages explaining background & what you did) due May 5
- * In-class presentations (5-10 minutes)
May 5 & 7

COURSE ADMIN / PROJECT

- * Presentations should be live (zoom) and include slides describing the background, your method/approach, and results
- * **It is essential** that every person in your group participates in both the project work and the presentation

COURSE ADMIN / PROJECT

- * Sources of neural data (among many more):
- * CRCNS: <https://crcns.org/data-sets>
- * Allen Inst.: <http://www.brain-map.org>
- * Study Forrest: <http://studyforrest.org/>

COURSE ADMIN / ?

- * Any further questions about any of these administrative things?

RECAP:

NONLINEAR METHODS


- * **Volterra series**
- * **Kernel regression** (*samples, not features!*)
- * **Artificial neural network** (*1 hidden layer*)

ARTIFICIAL NEURAL NETWORKS

- * We previously talked about **Perceptrons**, simple 1-layer networks with a step-function output nonlinearity
- * We looked at the Perceptron learning rule, which can be used to iteratively learn weights


ARTIFICIAL NEURAL NETWORKS

- * Generic 1-layer neural “network”
- * $g(z)$ is the output nonlinearity
- * Update rule? Just differentiate the loss function! (As long as $g(z)$ is differentiable)

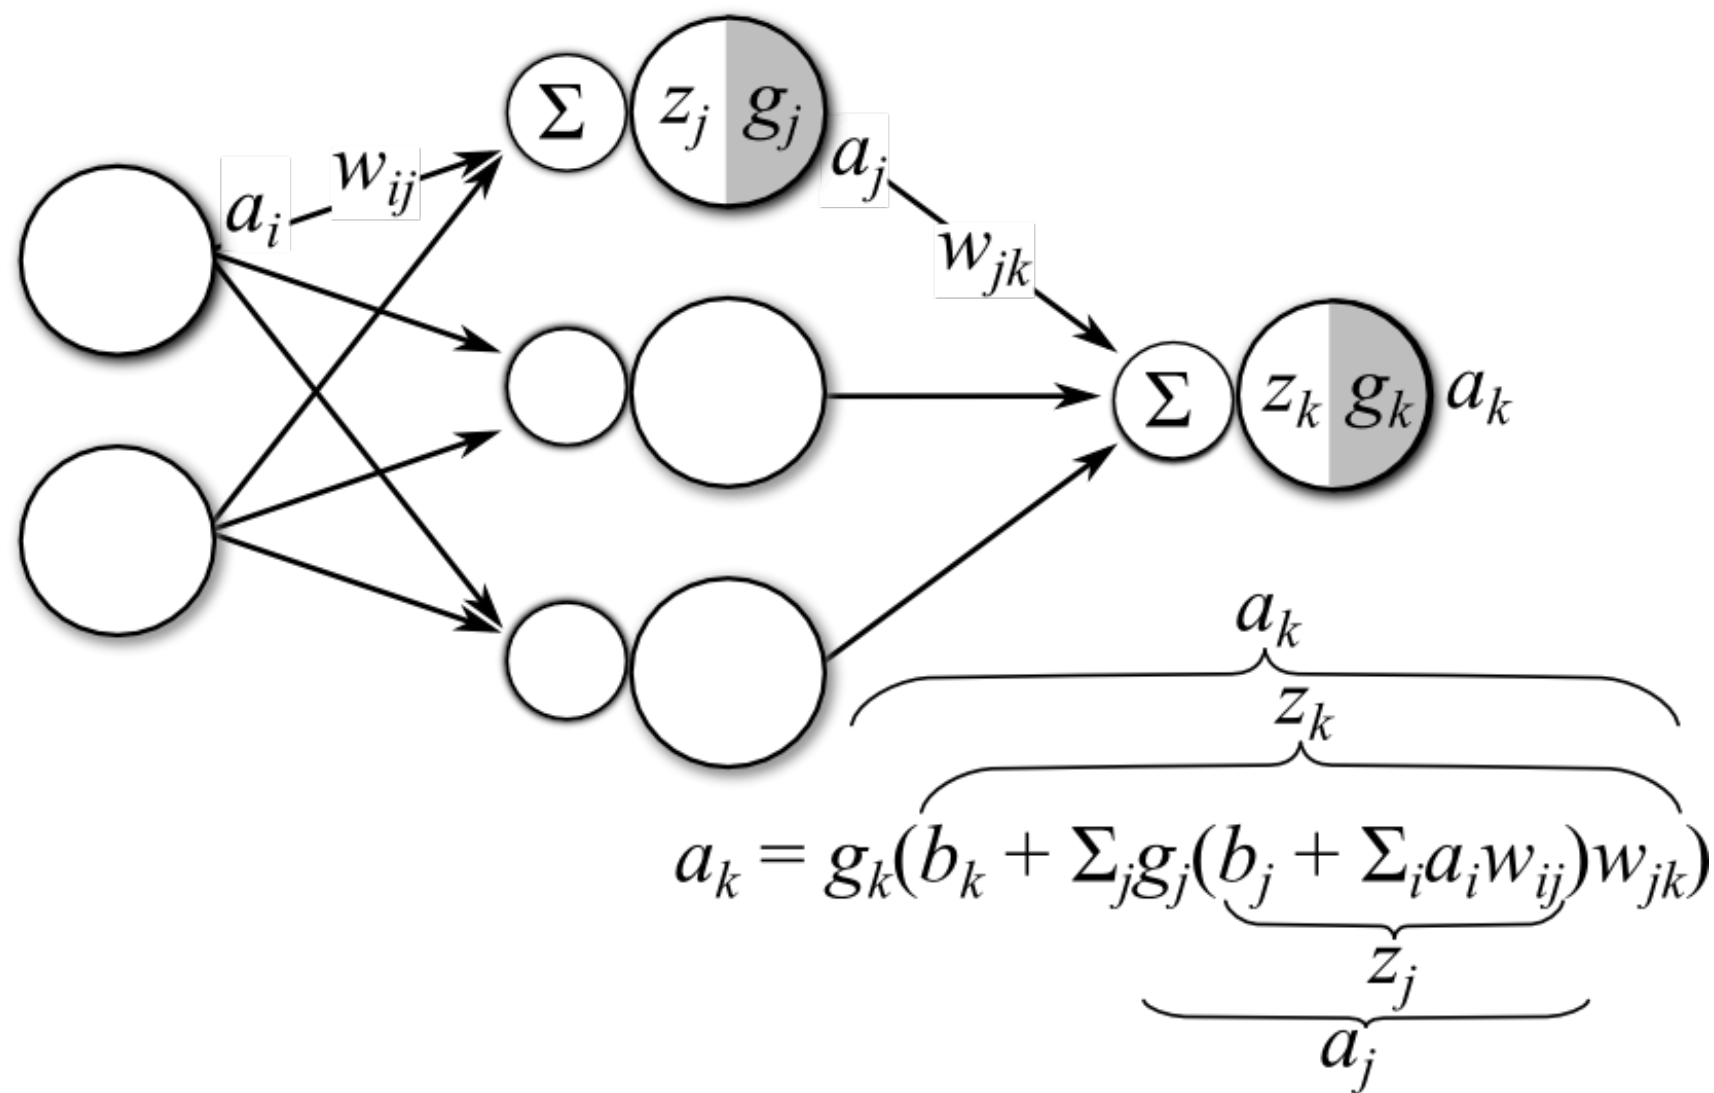


ARTIFICIAL NEURAL NETWORKS

- * Problem: this network is not much more “expressive” than the Perceptron
- * It still can't XOR
- * Solution: stack it! add a “hidden layer”



ARTIFICIAL NEURAL NETWORKS



w = weights


a = activations

z = total input

g = activ. fxn

b = bias

ARTIFICIAL NEURAL NETWORKS



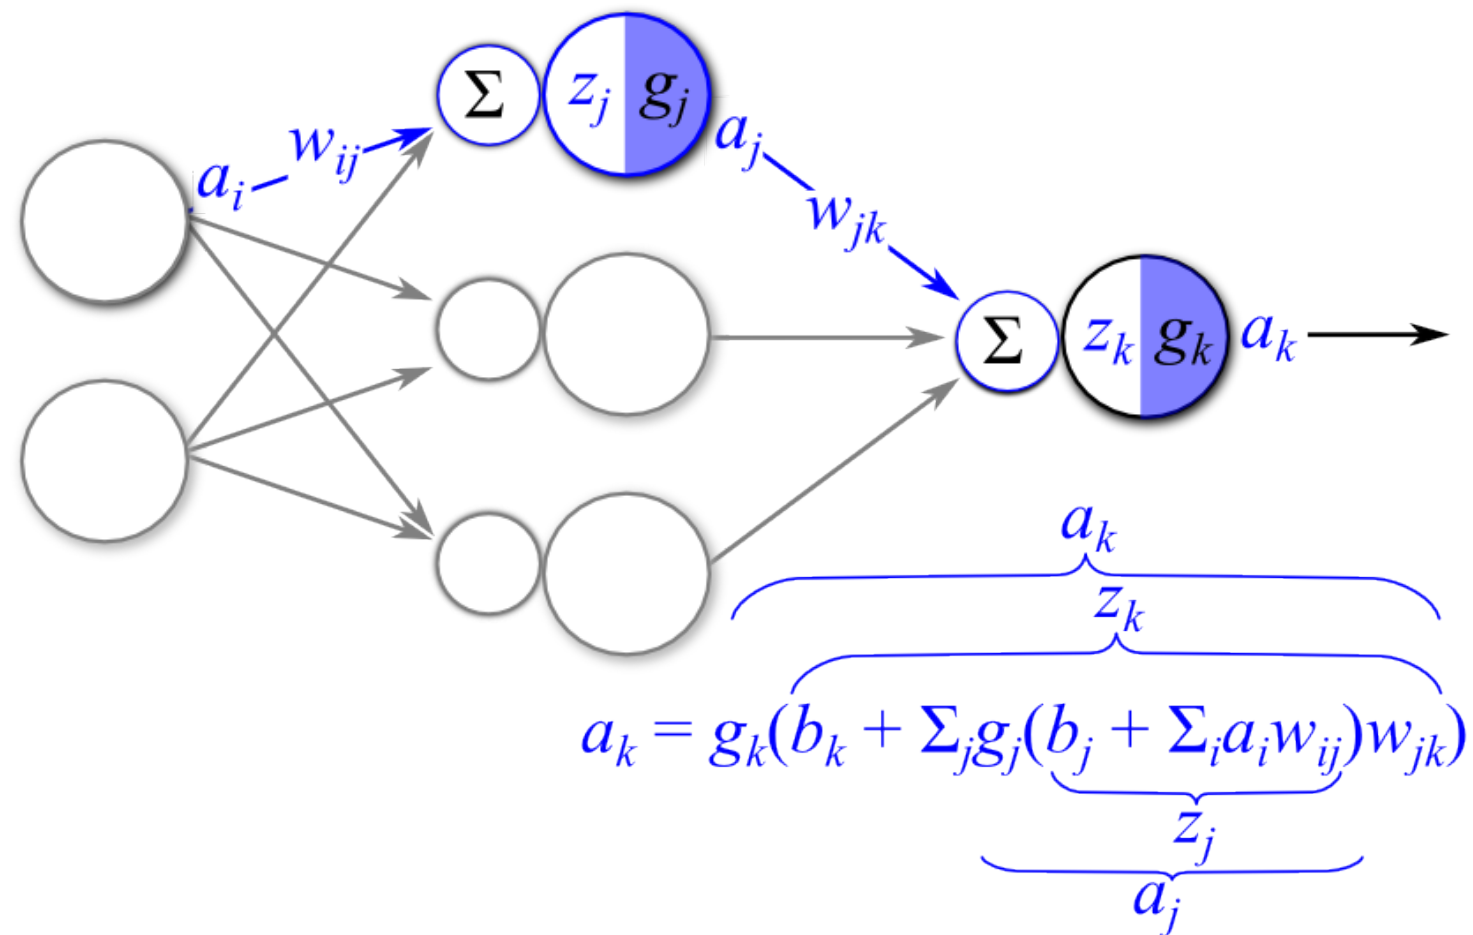
- * How do we train this beast?
- * Differentiate the loss function, same as ever!

GRADIENT BACKPROPAGATION

* (Derivation of backpropagation)


GRADIENT BACKPROPAGATION

I. Forward-propagate Input Signal




GRADIENT BACKPROPAGATION

II. Back-propagate Error Signals



GRADIENT BACKPROPAGATION

III. Calculate Parameter Gradients



GRADIENT BACKPROPAGATION

IV. Update Parameters

$$w_{ij} = w_{ij} - \eta(\partial E / \partial w_{ij})$$

$$w_{jk} = w_{jk} - \eta(\partial E / \partial w_{jk})$$

for learning rate η

GRADIENT BACKPROPAGATION

- * Backprop enables **efficient** computation of gradients in networks of **arbitrary depth** by caching intermediate values δ
- * It has probably been discovered in some form a few times in different fields
- * In this field we often credit Rumelhart, Hinton, & Williams (1986)

ARTIFICIAL NEURAL NETWORKS

- * What can networks with hidden layers do?
- * **Universal approximation theorem:** These networks can do *literally anything* (as long as there is a non-polynomial output nonlinearity)
- * See Cybenko (1989) paper attached to this lecture for proof

ARTIFICIAL NEURAL NETWORKS

- * But: it is not necessarily true that every function is *learnable* using backprop
- * Different network architectures (e.g. different numbers of hidden layers & units per layer) are good at learning different functions

NEXT TIME

- * Artificial neural networks for solving vision problems (e.g. categorization)