# LINKING ARTIFICIAL AND BIOLOGICAL NEURAL NETWORKS

Prof. Alexander Huth
4.7.2020

# COURSE ADMIN / PROJECT

* Proposal due **this** **Thursday** (**April 9**):

   * ~1-2 paragraphs describing what you plan to do & who is in your group. Email to huth@cs.utexas.edu before class

* Writeup (3-4 pages explaining background & what you did) due **May 5**

* In-class presentations (5-10 minutes) **May 5 & 7**

# RECAP

* **Deep** neural networks (i.e. with many hidden layers)

  * Trained using gradient **backpropagation**

  * Can do impressive things like **image classification** (e.g. **AlexNet**)

  * Using tricks like:

    * **localized** receptive fields

    * **convolution**

    * **downsampling** (i.e. max pooling)

# RECAP

* Subtle switch from **system identification** to **system construction**

* *System identification*: building functions to explicitly approximate **representations** used by a given computational system

* *System construction*: building functions with (ideally) the same **behavior** or **goal** as a given computational system

# RECAP: ALEXNET

* Does it solve the problem the same way that our brains do?

* How would we know?

# RECAP: ALEXNET

* Features learned
  by 1st layer of
  AlexNet

* Receptive fields
  of V1 neurons
  from a macaque

# CONSTRUCTION TO IDENTIFICATION

* What would it mean for a constructed system (e.g. AlexNet) to solve a problem **similarly** to how an existing system (e.g. a brain) does?

# CONSTRUCTION TO IDENTIFICATION

* Recall Marr's levels of description:

    * **Computational**: both systems solve the same task ✅

    * **Algorithmic/Representational**: both systems use the same algorithm or representations **? ? ?**

    * **Implementation**: both systems use the same implementation (neurons, etc.) ⛔

# CONSTRUCTION TO IDENTIFICATION

* Recall Marr's levels of description:

  * **Computational**: both systems solve the same task ✅

  * **Algorithmic/Representational**: both systems use the same algorithm or representations **???**

  * **Implementation**: both systems use the same implementation (neurons, etc.) ⛔

# CONSTRUCTION TO IDENTIFICATION

* What would it mean for a constructed system (e.g. AlexNet) to solve a problem **similarly** to how an existing system (e.g. a brain) does?

  * *It would mean that both use the same algorithm or representations to solve the problem*

  * How do we test whether a system uses some representation?

# CONSTRUCTION TO IDENTIFICATION

\* Key idea: *use the constructed system as a linearizing transform for system identification*

# EXTRACTING REPRESENTATIONS FROM A NEURAL NETWORK

* The **activations** at each layer of a deep neural network can serve as **"features"**

* The **functions** that compute the activations (from the input) thus serve as **linearizing transformations**

# EXTRACTING REPRESENTATIONS FROM A NEURAL NETWORK

* ("whiteboard" notes)

# CONSTRUCTION TO IDENTIFICATION

1. Collect dataset of brain responses *Y* using stimuli *X*

2. Train/find artificial system

3. Extract activations *a(X)* for the stimuli in the neuroscience experiment

4. Use activations as features in linearized encoding model: *Y = a(X) . b*

# CONSTRUCTION TO IDENTIFICATION

* Key questions:

  * Does this work at all?

    * If so, is this trivial?

  * Do activations from different layers predict different parts of the brain?

    * Does this correspond to known computational hierarchies?

    * If so, is *this* trivial?

# FMRI EXAMPLE

* Eickenberg, Gramfort, Varoquaux, & Thirion (2016) Seeing it all: convolutional network layers map the function of the human visual system. *Neuroimage*

# FMRI EXAMPLE

* fMRI data come from experiment we discussed previously (see *lecture 10*, Kay et al., Nature 2008)

* Two human subjects viewed 1000s of images while BOLD signals were recorded from visual cortex using fMRI

# FMRI EXAMPLE

* Used "Overfeat" convolutional neural network (similar architecture to AlexNet)

* Built model for each voxel using activations from each layer, then tested on held-out data



Input   Layer 1   Layer 5

Convolutional Net

Linear predictive models

Score maps ($R^2$)
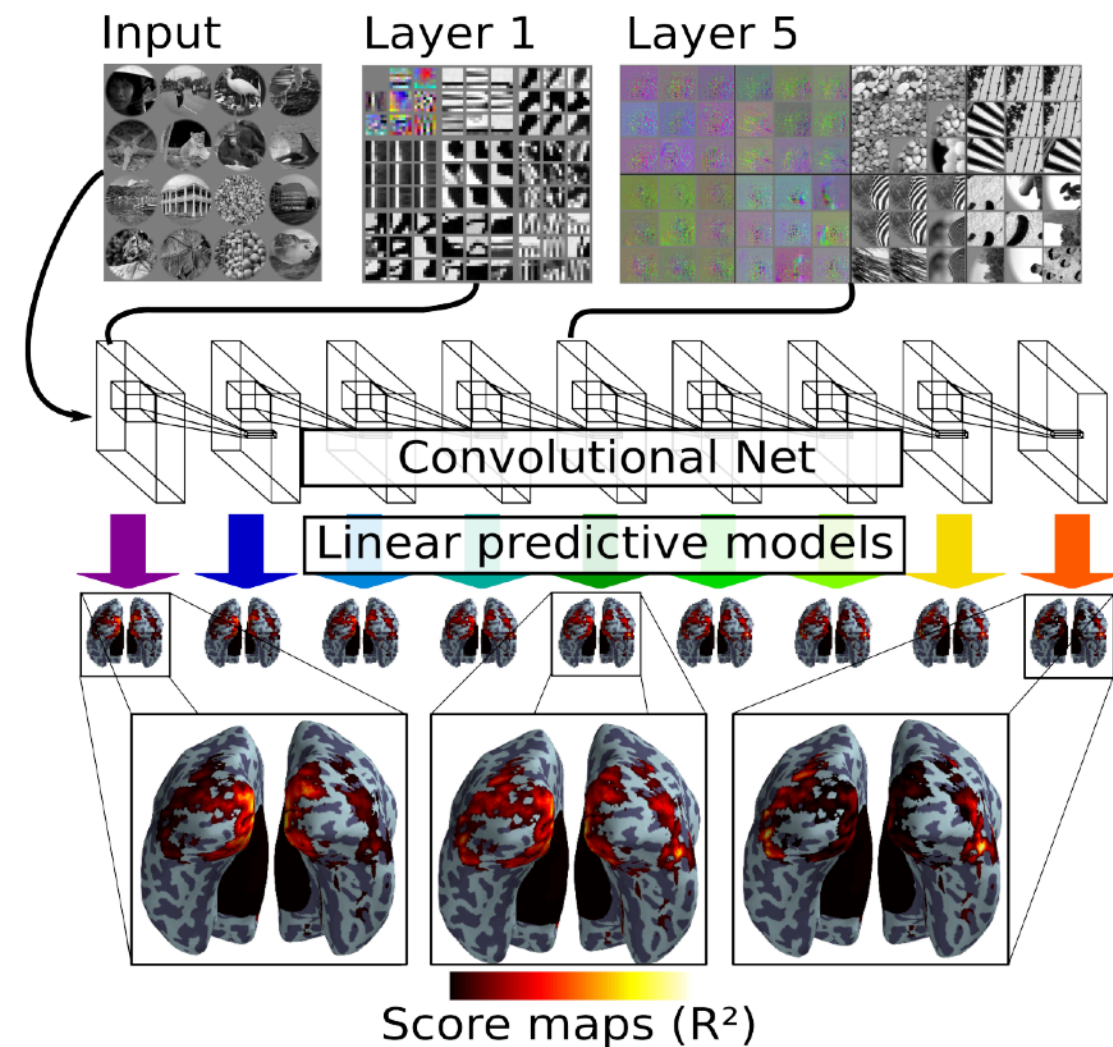
# VISUAL CORTEX

* **Connectivity** of areas in macaque visual cortex

* Felleman & Van Essen (1991)

# FMRI EXAMPLE



* Lower layers of the network predict "earlier" visual areas, higher layers predict "later" visual areas!
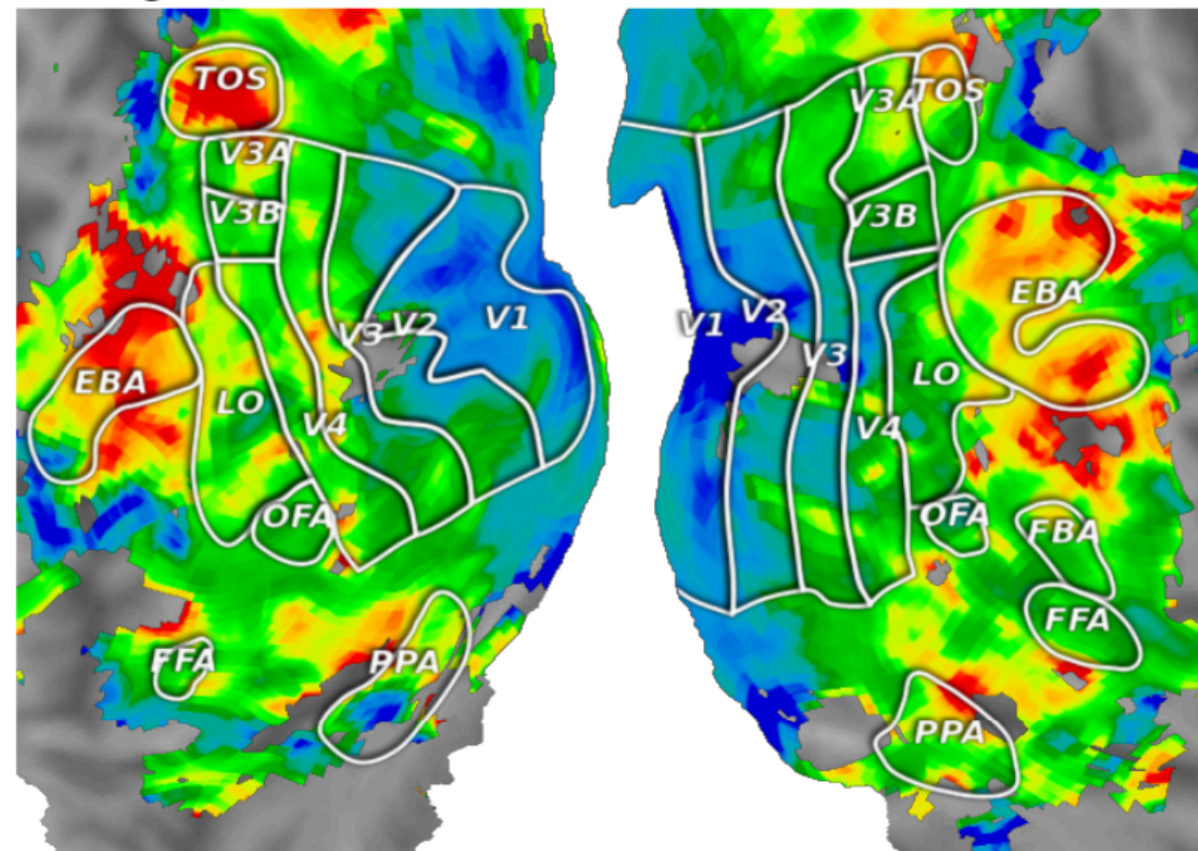
# FMRI EXAMPLE

# FMRI EXAMPLE

\* ***Fingerprint:*** summary model performance metric giving something like "best layer"



A Fingerprint summaries for Kay2008

B Fingerprint summaries for Huth2012

# CONSTRUCTION TO IDENTIFICATION

* Key questions:

  * Does this work at all? *YES*

    * If so, is this trivial? *HMM…*

  * Do activations from different layers predict different parts of the brain? *YES*

    * Does this correspond to known computational hierarchies? *YES*

    * If so, is *this* trivial? *HMM…*

# NEXT TIME

* Recurrent neural networks (etc.) for modeling language